

**Hamza Aziz**  
*AI Product Engineer @VoicePlug Inc.*

hamzaaziz822@gmail.com | +91 8447644188  
Sector-23, Sanjay Nagar, Ghaziabad-201002 (U. P.)  
*LinkedIn | GitHub | Portfolio | Medium*



## PROFESSIONAL SUMMARY

AI Product Engineer with 3 years of experience building computer vision, conversational AI, full-stack LLM applications, and applied MLOps systems for production environments. Delivered real-time facial recognition platforms on DGX A100, RTX 4090, and NVIDIA Jetson Orin Nano, reducing manual effort by up to 90% and achieving >99.99% recognition accuracy. Developed production voice-ordering and voice-to-SQL systems using LangGraph, Google Gemini, Rasa, FastAPI, Milvus, and retrieval-augmented workflows. Deploy production AI workloads with TensorRT, DeepStream, CI/CD automation, vector databases, and low-latency inference pipelines.

## SKILLS

- **Programming Languages:** Python, SQL, Bash
- **AI/ML:** PyTorch, TensorFlow, OpenCV, YOLO, InsightFace, Scikit-learn, NumPy, Pandas
- **Generative & Conversational AI:** LangChain, LangGraph, Rasa, Gemini, LLMs, NLP, Prompt Engineering, RAG, vLLM
- **Inference & Edge AI:** TensorRT, ONNX, DeepStream, NVIDIA Jetson, DGX A100
- **Backend & MLOps:** FastAPI, Next.js, React, TypeScript, Docker, CI/CD, Git, Linux
- **Databases & Messaging:** Milvus, PostgreSQL, MongoDB, Redis, Kafka, MinIO

## WORK EXPERIENCE

### AI Product Engineer

*VoicePlug Inc.* | September 2025 – Present

- Ship production-grade DriveThru voice ordering features for a multi-tenant Quick-Service Restaurant (QSR) platform deployed across 2 live restaurants and 4+ QSR brands
- Migrated core conversational workflows from legacy Rasa NLU to an agentic LangGraph + Google Gemini architecture, improving entity extraction, dialogue control, and tool-use orchestration
- Raised live client success metrics from ~25–30% to 60%+ by resolving intent-recognition gaps, entity ambiguity, menu-understanding failures, and fallback-handling issues
- Hardened POS-integrated deployments by improving exception handling, debug observability, fallback escalation, and restaurant-specific prompt templating, reducing recurring production issues by 6%
- Converted insights from 500+ real customer interactions into product improvements by partnering with product, support, and QA teams to reduce ordering friction and improve deployment readiness

### Assistant Engineer – AI

*Global InfoVentures Pvt. Ltd.* | August 2023 – August 2025

- Led development of 4 commercial-grade Facial Recognition Attendance Systems on DGX A100, RTX 4090, Jetson Orin Nano, achieving >99.99% accuracy
- Developed G6 Voice Assistant, converting natural language voice queries into precise SQL queries via Rasa, Gemini 2.0, and Milvus Vector DB; handled ambiguous queries with advanced intent disambiguation, achieving near-100% intent classification accuracy
- Mitigated domain shift in facial recognition using a CI/CD-style pipeline combining mobile and CCTV datasets — improved cross-domain recognition by 14%
- Integrated PostgreSQL for recognition logs and engineered an API sync module with the legacy SIM ERP Portal, enabling real-time attendance updates and SMS alerts for 2,000+ students

## EDUCATION

**Birla Institute of Technology and Science (BITS), Pilani – WILP**

2025 – 2027

*Master of Technology (Executive)*

- Artificial Intelligence and Machine Learning

**Dr. A. P. J. Abdul Kalam Technical University, ABES Institute of Technology**

*Bachelor of Technology*

2019 – 2023

- Computer Science and Engineering

## PROJECTS

**DriveThru Voice Ordering Modernization**

*VoicePlug Inc.*

- Shipped production features for a real-time, multi-tenant DriveThru voice ordering agent powered by FastAPI, Socket.IO, LangGraph ReAct workflows, Google Gemini 2.5 Flash fallback logic, and per-call token-cost accounting
- Refactored entity parsing, size normalization, and combo/multi-size guard logic to eliminate cart crashes and enforce semantically valid orders (e.g., wings overrides, mandatory “with Fries” confirmation) before persistence
- Hardened entity extraction against ASR disfluencies (partial names, fillers, mispronunciations) using fuzzy matching and added latency-masking filler scheduling that suppresses duplicate TTS when the real model response arrives in time
- Extended POS-compatible ordering workflows across 4 integrations (QUBeyond, Clover, Speedline, Deliverect) using Jinja2-templated restaurant prompts, Redis-backed session state, and MongoDB-backed order-event audit trails
- Raised production ordering reliability by tightening intent recognition, ambiguity handling, menu-constraint validation, and fallback/escalation behaviour across live QSR deployments

**G6 Voice Assistant – Natural Language to SQL Converter**

*Global InfoVentures Pvt. Ltd.*

- Integrated SpeechRecognition, Rasa, and Gemini 2.0 to convert user voice input into structured SQL queries, enabling seamless access to ERP data from modules like Admission, Employee, Fee, and Library
- Achieved ~100% intent classification accuracy using Rasa by training with 50+ voice query variations per module. Devised a secondary intent resolution via contextual follow-ups for ambiguous queries (example, "Hamza Aziz" as student or employee)
- Formulated 100+ field-level schema mappings per module for Gemini and integrated Milvus Vector DB-based suggestions, reducing LLM hallucinations and boosting SQL query precision by over 30%
- Crafted a voice-adaptive summary module that transformed 1000+ SQL outputs into natural language, enhancing end-user experience in production deployment
- Solved heterogeneity in field values (example, department / designation name variations) using vector-based field and value disambiguation, improving precision in SQL generation by over 40%

**Computer Vision Attendance Systems**

*Global InfoVentures Pvt. Ltd.*

- Deployed CCTV, PTZ, and edge-based facial recognition attendance systems using SCRFD, w600k\_r50, TensorRT, DeepStream, Milvus, DGX A100, RTX 4090, and Jetson Orin Nano
- Orchestrated 20+ live CCTV streams and automated face-vector refreshes from elevated-angle surveillance footage, improving real-world recognition accuracy from 85% to 99%
- Enabled low-latency, contactless attendance tracking across classroom, office, and surveillance environments, reducing manual roll calls, proxy attendance, and biometric contact dependency

## CERTIFICATES

- LangChain – Develop AI Agents with LangChain & LangGraph | Udemy - 2025
- Generative AI with Large Language Models | Coursera - 2024
- Building Video AI Applications at the Edge on Jetson Nano | NVIDIA Deep Learning Institute - 2024
- Getting Started with AI on Jetson Nano | NVIDIA Deep Learning Institute - 2024
- Artificial Intelligence and Machine Learning | Global InfoVentures Pvt. Ltd. - 2023
- Getting Started with Deep Learning | NVIDIA Deep Learning Institute – 2022